

9. Operating Systems.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Operating Systems

- Purpose: to provide high level facilities beyond the capabilities of the hardware
- To provide security for data
- An operating system is really a collection of a large number of services.
 - Some are essen
 - Others are avai
- Examples: Windows, Unix, mainfra
- In general modern operating systems are big
 - E.g. Windows 50,000,000 source lines of code.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

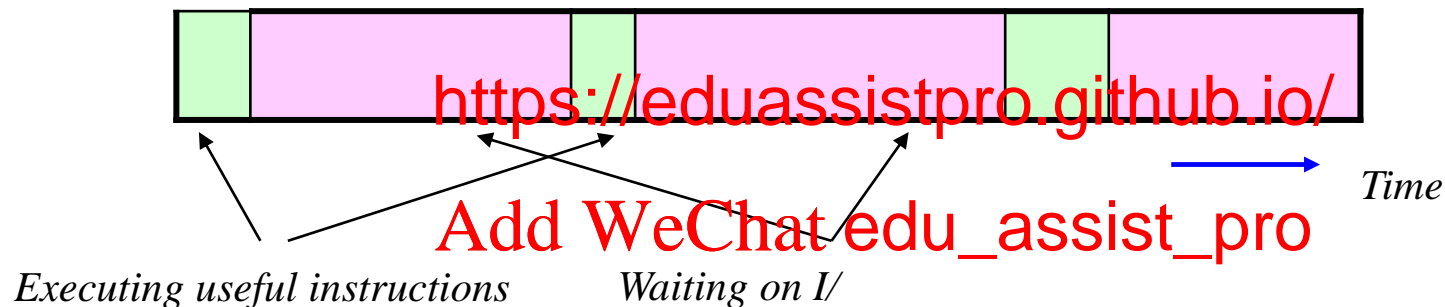
request (usually via an

al-time systems.

Historical Motivation: Buffered I/O

- The problem: calculations are fast, I/O is usually slow
 - E.g. reading a character of input is much slower than executing a statement like $x = 2*3$
 - Therefore, if we just have a bare machine running a program, most of the time will be spent waiting on I/O, with little useful work getting done

Assignment Project Exam Help



- An idle CPU is a wasted resource. Solution: find it something else to do!
- *Example:* when a computer prints a document, it doesn't just freeze up until the paper is sitting there with ink on it! It continues to perform other tasks.


A Solution: Concurrent Processes

- Introduce the idea of **processes**
- A process is a running program, with its evolving state and dynamic behaviour
- It seems like several processes can run at the same time on a single CPU. This is called **multi-tasking**.
 - In reality, the OS uses **timer interrupts** and **time slices** to implement multitasking.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



The OS Controls Processes

- There may be many independent processes running concurrently. But:
 - When waiting for I/O how would an application know which other application to jump to?
 - How does an ordinary application even know what programs are running?
- If we want multitasking:
 - What about I/O devices which multiple processes may need (screen, keyboard, printer etc.)?
 - How do we stop two independent processes from over-writing each other's memory?
- Solution: the **Operating System**
 - Maintains a table
 - Mediates every process break, deciding
 - Controls which process gets to access wh any given time
 - Controls access to physical memory with the help of a hardware device capable of translating outgoing addresses (a **memory management unit** or **MMU**).
- Process control is the most fundamental facility provided by operating systems
- It's implemented by the most central code in the OS, called the **kernel**.

Multitasking on one CPU

- At any moment in time, the CPU is executing just one instruction!
 - This belongs either to one of the processes, or to the operating system
 - One process is running, the others are waiting
- CPU switches execution between the processes rapidly (~ 100 times per second)
- User's point of view
 - Computer responds smoothly to keystrokes and mouse motions
 - In fact CPU is run
- Computer's point of view
 - Rapidly alternates attention among many p
- The Operating System maintains a list of g at the same time
 - What does “at the same” time mean? Depends on time scale:
 - *At any given instant*, hardware is executing one instruction belonging to one program
 - *Over a longer period of time*, as interrupts keep transferring control between programs, all programs get some execution time.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Processes

- **A *program* is a document:** a text file containing a source program, or a binary file containing a machine language program.
 - A program just sits there, like a book.
- **A *process* is a running program.** May not be executing all the time but soon an interrupt will cause the computer to resume its instructions
- At any moment it is in one of the following states:
 - **Running:** at the time the process is executing an instruction belonging to it. The process is currently in the computer is
 - **Blocked:** the process is waiting for I/O
 - **Ready:** the process is not waiting for I/O, so it could be running, but the CPU is doing something else

Concurrency and Parallelism

- Related but distinct concepts that are often confused
- Several processes run **concurrently** when all make progress on user's time scale
 - Can involve parallelism, but usually just the CPU switching rapidly among the processes.
- Processes run in **parallel** when they execute simultaneously.
 - Requires multiple processors, used for high performance applications
 - Many modern microprocessors are essentially multiple CPUs on one chip and can execute multiple processes concurrently (not truly in parallel).
 - An individual core can execute only one process at a time.
- In order to get smooth concurrency on one processor, the OS must switch rapidly and regularly among the processes even if the running process has no requested I/O.
 - If we don't, computer won't respond smoothly to user requests.
 - Even with multiple CPUs, there are usually not enough for all processes so, each must also switch processes rapidly and regularly.

Co-operative & Pre-emptive Scheduling

- Simplest approach is **co-operative scheduling**:
 - All applications must be written so as to voluntarily give up control to the OS at regular periods.
 - The relinquish operation is called a **process break** and accesses via a TRAP instruction with appropriate parameter.
 - If an application fails to perform a process break, it gets all the machine's resources. All other processes are **starved**.
 - This is a poor solution.
 - A process break is not terminate it.
- Much better is **pre-emptive scheduling**.
 - Hardware and OS work together to enforce a priority-regular process breaks.
 - The hardware must have the ability to signal to the CPU at any given time that it must stop its current sequence of instruction cycles.
 - This is provided via special control inputs to the CPU called interrupt lines.
 - Needs extra hardware such as timer.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Physical I/O

- I/O devices have control/status registers as well as data registers.
 - These locations may appear in the normal physical address space and are accessed by usual read and write operations (**memory mapped**).
 - Alternatively they may be accessible by special I/O instructions.
 - Some systems use both approaches (e.g. a PC graphics card)
 - Storing data in a data register for an output device sends the word to the device
 - Reading data from the data from the device.
- Some I/O devices can be installed once; others can only service one at a time.
- Often I/O devices are slow so process request to wait.
- A process may wait by looping and **polling** it) or it may be notified when request completes. In latter case it may be suspended by OS (**synchronous I/O**) or it may carry on running (**asynchronous I/O**).
- A suitable OS can relieve applications of responsibility for directly controlling I/O devices and can ensure that each device is shared appropriately between applications.
- When I/O is under OS control, application will normally access a device via high-level commands. These are translated into physical access by an OS **device driver**.

Interrupts

- A **hardware interrupt** is a feature that allows an external hardware device to indicate to a CPU that it should perform a process break.
- When an interrupt occurs the CPU is directed to an **interrupt handler routine** which has been prepared in advance and which it begins executing.
 - As with a subprogram, the state of the CPU must be saved so that the CPU can return to the interrupted program ready to resume execution.
 - Except in the very simplest systems, the interrupt handler routine is in the OS and the OS decides what to do next.
 - An interrupt can occur at any time but the CPU *always finishes the current instruction before jumping to the handler.*
- Interrupts may be generated at arbitrary intervals by I/O devices that need service or regularly by a **system timer** (see next).

Interrupt Handling

- In a pre-emptively scheduled system a special hardware timer is used to generate interrupts at regular intervals (e.g. every 100ms).
 - This invokes a special interrupt routine inside the OS called the **process scheduler**.
 - The process scheduler uses a **scheduling algorithm** to choose which process goes next (this may include some concept of **priority**).
 - No process gets more time than the **time slice**.
- Interrupts are not initiated by a process and occur at the end of any instruction.
- An interrupt is only supposed to suspend a process not terminate it.
- In order to resume a process from any arbitrary point, all registers in the programmer's model must be saved when an interrupt handler is invoked. **Why just the programmer's model?**